

```
1.
library(gdata)
diamondData<-read.xls("Diamond_Data.xls", perl = "c:/Perl64/bin/perl.exe")
```

```
head(diamondData)
```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Diamond_data/ ↗

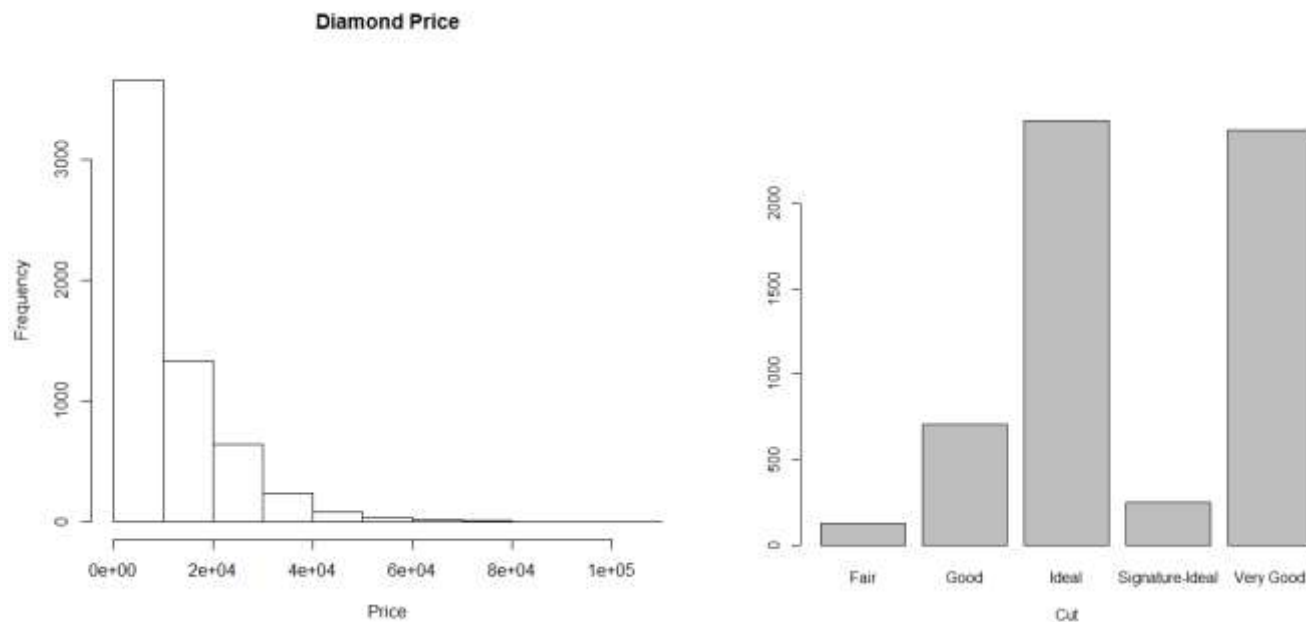
```
> head(diamondData)
```

	ID	Carat.weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
1	1	1.10	Ideal	H	SI1	VG	EX	GIA	5169
2	2	0.83	Ideal	H	VS1	ID	ID	AGSL	3470
3	3	0.85	Ideal	H	SI1	EX	EX	GIA	3183
4	4	0.91	Ideal	E	SI1	VG	VG	GIA	4370
5	5	0.83	Ideal	G	SI1	EX	EX	GIA	3171
6	6	1.53	Ideal	E	SI1	ID	ID	AGSL	12791

```
#To install Rule Fit
platform = "windows"
rfhome = "C:/Program Files/R/RFHOMe"
source("C:/Program Files/R/RFHOMe/rulefit.r")
install.packages("akima", lib=rfhome)
library(akima, lib.loc=rfhome)
#end of rulefit install
```

1.a Show the distribution of the “cut” and “Price” attributes.

```
hist(diamondData$Price, main="Diamond Price", xlab="Price")
plot(diamondData$Cut , xlab="Cut")
```



```
1.b
per <- floor(nrow(diamondData)*5/6)
subs <- sample(nrow(diamondData),per)
train <- diamondData[subs,]
x <- train[,2:8]
y <- train[,9]
cat.var <- c("Cut", "Color", "Clarity", "Polish", "Symmetry", "Report")
rfit <- rulefit(x,y,rfmode="regress",cat.vars=cat.var,test.reps=10,test.fract=0.1)
```

rfmodinfo(rfit)

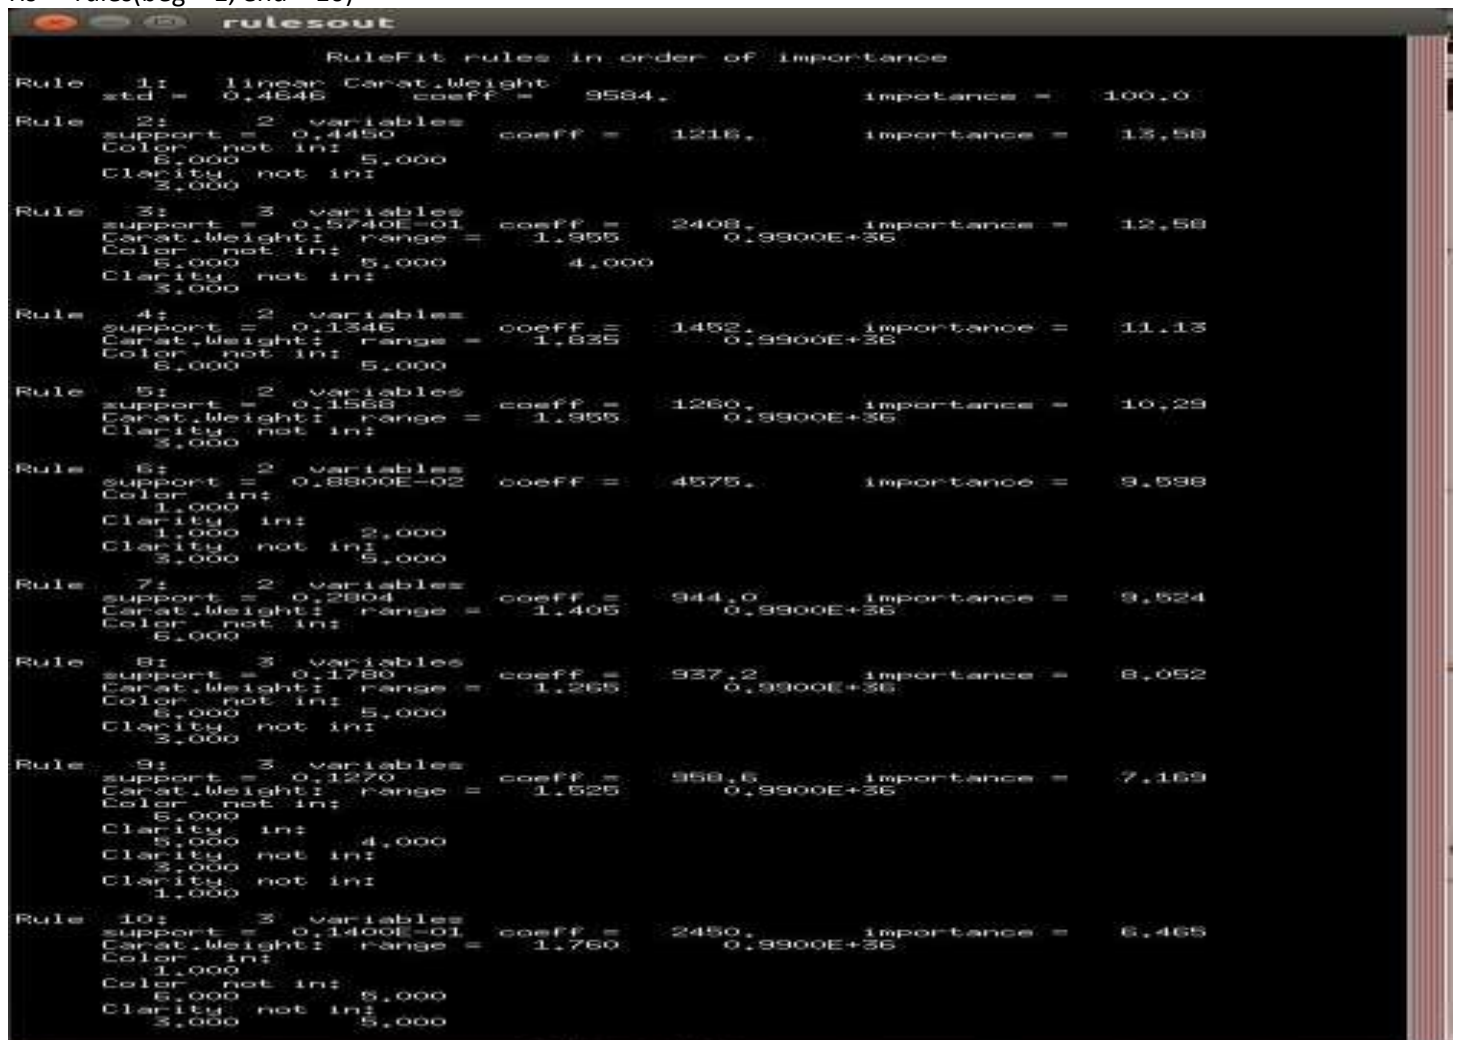
```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Diamond_data/
> rfmodinfo(rfit)
rulefit(x = x, y = y, cat.vars = cat.var, rfmode = "regress",
  test.reps = 10, test.fract = 0.1)
RuleFit model 12/02/2015 2:45p
estimated: criterion value      +/-      # terms      569
          0.1699E+07          0.1498E+06
Parameters:
  cat.vars = Cut Color Clarity Polish Symmetry Report
  not.used =
  xmiss = 9e+30
  rfmode = regress
  sparse = 1
  test.reps = 10
  test.fract = 0.1
  mod.sel = 2
  model.type = both
  tree.size = 4
  max.rules = 2000
  max.trms = 500
  trim.qnt1 = 0.025
  samp.fract = 0.1557635
  inter.supp = 3
  memory.par = 0.01
  conv.thr = 0.001
```

Average absolute error: 0.1498E+06

Number of terms: 569

1.c

rls <- rules(beg = 1, end = 10)




Rule	support	variables	coeff	importance
Rule 1: linear	0.4648	Carat,Weight	9584.	100.0
Rule 2: 2 variables	0.4450	Color not in: 5.000 Clarity not in: 5.000	1216.	13.58
Rule 3: 3 variables	0.5740E-01	Carat,Weight: range = 1.955 Color not in: 5.000 Clarity not in: 4.000	2408. 0.9900E+36	12.58
Rule 4: 2 variables	0.1346	Carat,Weight: range = 1.835 Color not in: 5.000	1452. 0.9900E+36	11.13
Rule 5: 2 variables	0.1568	Carat,Weight: range = 1.955 Clarity not in: 5.000	1260. 0.9900E+36	10.29
Rule 6: 2 variables	0.8800E-02	Color in: 1.000 Clarity in: 1.000 Clarity not in: 2.000 Clarity not in: 3.000	4575.	9.598
Rule 7: 2 variables	0.2804	Carat,Weight: range = 1.405 Color not in: 5.000	944.0 0.9900E+36	9.524
Rule 8: 3 variables	0.1780	Carat,Weight: range = 1.265 Color not in: 5.000 Clarity not in: 3.000	937.2 0.9900E+36	8.052
Rule 9: 3 variables	0.1270	Carat,Weight: range = 1.525 Color not in: 5.000 Clarity in: 5.000 Clarity not in: 4.000 Clarity not in: 3.000 Clarity not in: 1.000	958.6 0.9900E+36	7.169
Rule 10: 3 variables	0.1400E-01	Carat,Weight: range = 1.760 Color in: 1.000 Color not in: 5.000 Clarity not in: 5.000 Clarity not in: 3.000	2450. 0.9900E+36	6.465

After observing the above output,

- i. There was a liner term in the model 'carat.weight'. The coefficient is 9584 which suggests increase in the weight. The price will increase by \$9584
- ii. Rule 3 is based on 3 variables and has importance of 12.58. The rule indicates that if the carat.weight is in the range of 1.955 and color is not in 6, 5, 4 and clarity is not 3 then price will increase by 2408\$

1.d.

```
v <- varimp(range = 1:nrow(diamondData[-subs,]),x=diamondData[-subs,][,2:8],wt=rep(1,nrow(diamondData[-subs,])),rth=0,col='grey', donames=T, las=2)
```



```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Diamond_data/
> v
$imp
[1] 100.0000000 43.9227801 43.7858091 0.9919589 0.0000000 0.0000000 0.0000000

$sord
[1] 1 3 4 6 2 5 7
```

As seen from the plot the top three most important variables in determining the diamond's price are: Carat.Weight, Color and Clarity

1.e.

```
AvgAbsError <- function(predictedValues, actualValues){
  sum=0
  for(i in 1:length(predictedValues)){
    sum = sum + actualValues[i]-predictedValues[i]
  }
  return(sum/length(predictedValues))
}
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Diamond_data/
> AvgAbsError <- function(predictedValues, actualValues){
+   sum=0
+   for(i in 1:length(predictedValues)){
+     sum = sum + actualValues[i]-predictedValues[i]
+   }
+   return(sum/length(predictedValues))
+ }
```


```
testData <- diamondData[-subs,2:8]
actualValues <- diamondData[-subs,9]
predictedValues <- rfpred(testData)
averageAbsError <- AvgAbsError(predictedValues,actualValues)
averageAbsError
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Diamond_data/
> testData <- diamondData[-subs,2:8]
> actualValues <- diamondData[-subs,9]
> predictedValues <- rfpred(testData)
> averageAbsError <- AvgAbsError(predictedValues,actualValues)
> averageAbsError
[1] 26.59627
```

Average Absolute error: 26.59627

1.f.

```
x<- train[,2:9]
X1<-as(x,"data.frame")
X2<-as(X1,"data.frame")
install.packages("tree")
library(tree)
dtr<-tree(Price~.,X1)
plot(dtr)
text(dtr, cex=.8,pretty=0)
```



```
install.packages("rpart")
library(rpart)
```

```
ftO<-rpart(Price ~ .,data=X2,method="class",cp=0.0001)
printcp(ftO)
plotcp(ftO)
```

```
> printcp(fto)


Classification trees:
rpart(formula = Price ~ ., data = x2, method = "class", cp = 1e-04)

Variables actually used in tree construction:
[1] Carat.weight clarity color cut Polish Symmetry

Root node error: 4993/5000 = 0.9986

n= 5000
```

	CP	nsplit	rel error	xerror	xstd
1	0.00100140	0	1.00000	1.00120	0.00020026
2	0.00080112	1	0.99900	1.00080	0.00034679
3	0.00075105	5	0.99579	0.99940	0.00063271
4	0.00060084	13	0.98979	0.99880	0.00072118
5	0.00050070	27	0.98137	0.99539	0.00109368
6	0.00048640	29	0.98037	0.99519	0.00111165
7	0.00040056	36	0.97697	0.99139	0.00140910
8	0.00037195	162	0.92650	0.99079	0.00145031
9	0.00030042	169	0.92389	0.99059	0.00146379
10	0.00020028	177	0.92149	0.99039	0.00147713
11	0.00010014	407	0.87543	0.98778	0.00164029
12	0.00010000	409	0.87523	0.98778	0.00164029



```
fto.cpt<-ftO$cptable
min(fto.cpt[,"xerror"])
```

```
Console D:/Courses/Pattern Recognition/a
> min(fto.cpt[,"xerror"])
[1] 0.9877829
```

```
row_min<-which(fto.cpt[,"xerror"]==min(fto.cpt[,"xerror"]))
fto.cpt[row_min,"CP"]
```

```
Console D:/Courses/Pattern Recognition/a
> fto.cpt[row_min,"CP"]
      11      12
0.0001001402 0.0001000000
```

```
best_cp<-fto.cpt[row_min,"CP"]
dtree<-prune(ftO, best_cp[1])
```

Error: 0.9877829

2.


2.a

```
letters<-read.table("az-5000.txt", header=FALSE)
#not selecting first column
letters<-letters[2:5001,2:19]
letters.3means<-kmeans(letters,centers=26)
totwithinss=c()
for(i in 2:26){
  totwithinss[i]=(kmeans(letters,i)$tot.withinss)/i
}
```


K	J(C)
2	2486.86724
3	1462.92239
4	1001.30631
5	753.04704
6	582.54581
7	456.87470
8	377.80740
9	318.00262
10	280.63136
11	240.86812
12	218.09774
13	192.81951
14	173.09759
15	155.23682
16	142.99360
17	129.09894
18	118.23988
19	111.03241
20	103.47979
21	95.24959
22	88.96936
23	84.98184
24	80.69763
25	75.83029
26	70.09717

2.b

```
plot(totwithinss,1:26, type = "b", xlab = "Goodness of fit", ylab = "K")
```



```
plot(totwithinss[15:26],15:26, type = "b", xlab = "Goodness of fit", ylab = "K")
```




The above plot shows the goodness of fit for $K=15$ to 26 . As seen from the plot above, As the number of clusters increases, the value of withinss decreases. In the above plot, at $k = 20$, there is a step. This indicates the 20th letter 'T' might suggest the number of natural clusters.

3

3.a.

```
data <- read.table("az-5000.txt", header=TRUE)
azletters<-read.table("az-5000.txt", header=TRUE)
azletters<-azletters[,-1]
azletters.3means<-kmeans(azletters,centers=26, nstart = 22)
hc <- hclust(dist(azletters.3means$centers), method = "average")
plot(hc)
```



3.b.

```
characts <- data[,1]
charactMatrix <- table(characts,azletters.3means$cluster)
maxCharsArray <- array(26)
for(i in 1:26){
  maxCharsArray[i] <- letters[which.max(charactMatrix[,i])]
}
```



```

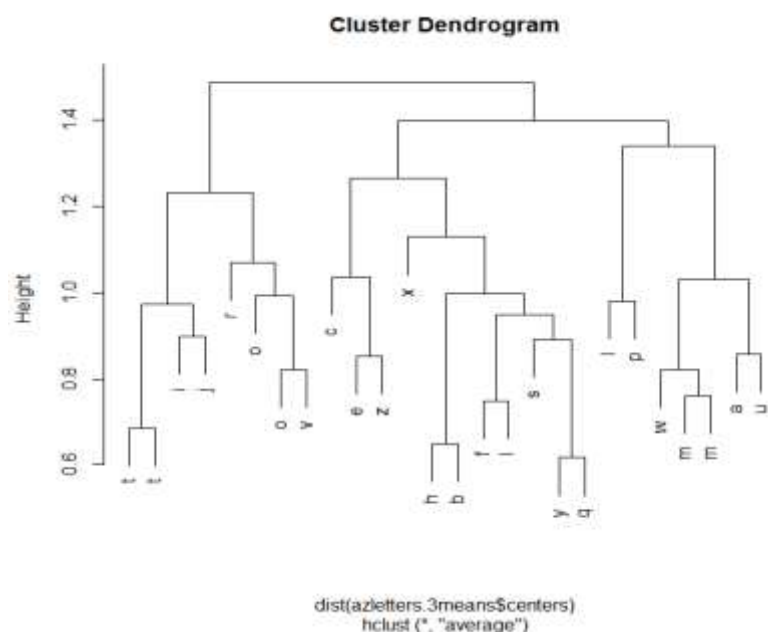
Console D:/Courses/Pattern Recognition/assignments/HW4/data/ >
> charactMatrix
characts 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26
a 0 1 86 0 0 0 0 0 0 0 0 0 0 11 8 0 0 2 4 1 40 1 1 13 0 10
b 1 8 0 9 0 0 0 0 145 5 4 0 0 0 2 2 0 0 0 0 0 0 1 0 0 7 0
c 0 0 1 7 0 0 0 0 0 0 0 0 0 4 0 0 0 0 177 3 3 1 0 3 0 0
d 1 9 26 0 1 4 0 1 0 13 0 0 0 0 87 1 1 0 0 2 26 1 2 10 0 2
e 1 0 0 179 0 0 0 0 0 2 0 0 0 0 0 0 0 0 5 0 2 0 0 2 2 2
f 4 0 0 4 0 0 0 0 0 26 5 0 0 0 99 0 0 15 0 1 3 1 25 0 0 3 1
g 102 0 1 0 0 0 0 0 33 0 0 0 0 34 0 1 1 0 0 1 2 3 14 0 0 0
h 0 146 4 9 0 0 0 2 0 0 0 0 0 0 4 0 0 2 1 7 1 0 0 2 0 0
i 0 0 0 0 1 62 0 0 0 2 0 0 111 0 0 0 2 3 0 9 0 0 1 1 0
j 2 0 0 0 0 14 0 0 13 0 1 0 151 3 0 1 0 0 0 0 1 8 0 0 0 0
k 3 109 0 3 0 0 0 0 0 0 0 0 3 0 2 22 2 2 2 0 4 2 2 15 1 0 1
l 2 2 0 12 0 0 0 0 0 30 0 0 0 1 0 0 1 0 2 150 0 0 0 0 1 1
m 0 0 0 0 0 0 0 0 0 0 0 136 0 0 1 0 0 4 0 0 1 0 1 49 0 0
n 0 18 0 0 0 0 0 0 0 2 2 11 1 0 100 0 4 2 0 0 0 0 6 18 0 0
o 2 0 1 0 74 0 106 0 1 3 2 0 0 2 1 3 0 1 5 0 0 2 0 3 2 3
p 2 1 1 0 0 0 0 0 0 8 9 0 0 0 1 0 177 0 0 1 0 4 2 2 0 0
q 24 1 3 0 0 0 0 0 1 0 0 0 0 17 1 0 2 0 0 0 0 0 148 0 0 0
r 1 1 0 5 0 3 0 0 0 9 165 0 0 0 11 3 0 0 0 0 10 2 6 1 0
s 3 0 0 0 0 0 0 187 1 1 0 0 3 0 0 0 0 4 0 1 0 7 0 0
t 0 0 1 3 1 1 3 0 0 3 5 0 1 7 0 2 0 0 0 3 3 69 0 2 87 0
u 0 20 0 0 0 0 0 0 0 0 0 1 0 0 143 26 0 4 1 0 0 0 3 21 1 0
v 0 0 0 0 2 0 0 0 0 0 8 0 0 0 4 177 0 2 0 0 0 4 1 2 0 0
w 0 0 0 0 0 0 0 0 0 0 0 1 0 0 2 10 0 180 0 0 0 0 0 5 0 0
x 4 3 53 3 3 0 2 0 0 8 0 0 0 1 3 20 1 3 1 0 1 12 10 3 0 52
y 148 0 0 0 0 0 0 0 10 0 2 0 0 20 0 0 0 0 0 1 1 0 6 0 0 3
z 1 0 0 0 0 0 0 0 0 0 0 0 0 9 0 0 0 0 0 0 115 1 0 0 34 0

```

```

hc$labels <- maxCharsArray
plot(hc)

```



The most common letters in the clusters are 'n' and 'u'.

3.c.

- From the dendrogram it can be observed that, deeper the location of an item in the dendrogram, higher is its frequency of occurrence in the clusters.
- Similar letters are present together in the dendrogram

3.d.

The missing letters are: d, n, g, k.

The missing letters should be assigned to the cluster in which they have maximum occurrence.

Letter	Cluster
n	15
G	1
d	15
k	2

4.
require(arules)
4.a.
ratingsAsBasket <- read.transactions("ratingsAsBasket.txt", format = "basket")
summary(ratingsAsBasket)

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> summary(ratingsAsBasket)
transactions as itemMatrix in sparse format with
10000 rows (elements/itemsets/transactions) and
15500 columns (items) and a density of 0.009911529

most frequent items:
M.4712.R.High M.3749.R.High M.5407.R.High M.4275.R.High M.538.R.High (other)
4729 4610 4162 4152 4010 1514624

1208 1209 1212 1216 1219 1225 1230 1245 1255 1272 1285 1
1 1 1 1 1 1 1 1 1 1 1 1
1666 1709 1852 1945 1972 2003 2027 2087 2106 2267 2289
1 1 1 1 1 1 1 1 1 1 1 1

Min. 1st Qu. Median Mean 3rd Qu. Max.
20.0 47.0 92.0 153.6 183.0 2289.0

includes extended item information - examples:
Labels
1 M.1.R.High
2 M.1.R.Low
3 M.1.R.Med

```

Number of baskets: 10000
Most frequent item: M.4712.R.High
Title: The Matrix (4712)
Rating: High
Frequency: 4729
Number of movies rated by one rater:
Minimum: 20.0
Maximum: 2289.0
Average: 92.0

4.b.
apr <- apriori(ratingsAsBasket)

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> apr <- apriori(ratingsAsBasket)
Apriori

Parameter specification:
confidence minval smax arem aval originalsupport support minlen maxlen target ext
0.8 0.1 1 none FALSE TRUE 0.1 1 10 rules FALSE

Algorithmic control:
filter tree heap memopt load sort verbose
0.1 TRUE TRUE FALSE TRUE 2 TRUE

Absolute minimum support count: 1000

set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[15500 item(s), 10000 transaction(s)] done [0.26s].
sorting and recoding items ... [253 item(s)] done [0.03s].
creating transaction tree ... done [0.01s].
checking subsets of size 1 2 3 4 5 6 done [0.10s].
writing ... [571 rule(s)] done [0.00s].
creating s4 object ... done [0.01s].

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> apr
set of 571 rules

```

Summary(apr)

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> summary(apr)
set of 571 rules

rule length distribution (lhs + rhs):sizes
  2   3   4   5
  3 170 357  41

    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 2.000  3.000   4.000   3.764  4.000   5.000

summary of quality measures:
      support      confidence      lift
Min.   :0.1000  Min.   :0.8000  Min.   :1.692
1st Qu.:0.1034  1st Qu.:0.8115  1st Qu.:1.988
Median :0.1080  Median :0.8222  Median :2.113
Mean   :0.1104  Mean   :0.8258  Mean   :2.171
3rd Qu.:0.1142  3rd Qu.:0.8385  3rd Qu.:2.297
Max.   :0.1565  Max.   :0.8806  Max.   :3.143

mining info:
      data ntransactions support confidence
ratingsAsBasket      10000      0.1      0.8
```

Top 10 rules: With respect to 'lift' measure

```
inspect(head(sort(apr, by = "lift"),10))
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> inspect(head(sort(apr, by = "lift"),10))
  lhs                                     rhs      support confidence lift
199 {M. 2250. R. High,M. 2936. R. High,M. 647. R. High} => {M. 646. R. High} 0.1025 0.8464079 3.142993
205 {M. 2526. R. High,M. 2749. R. High,M. 647. R. High} => {M. 646. R. High} 0.1007 0.8440905 3.134387
215 {M. 2250. R. High,M. 4275. R. High,M. 647. R. High} => {M. 646. R. High} 0.1157 0.8390138 3.115536
224 {M. 2526. R. High,M. 4275. R. High,M. 647. R. High} => {M. 646. R. High} 0.1119 0.8369484 3.107866
208 {M. 2250. R. High,M. 2526. R. High,M. 647. R. High} => {M. 646. R. High} 0.1158 0.8324946 3.091328
203 {M. 2250. R. High,M. 2749. R. High,M. 647. R. High} => {M. 646. R. High} 0.1006 0.8293487 3.079646
235 {M. 4275. R. High,M. 4712. R. High,M. 647. R. High} => {M. 646. R. High} 0.1112 0.8261516 3.067774
218 {M. 2250. R. High,M. 4712. R. High,M. 647. R. High} => {M. 646. R. High} 0.1130 0.8242159 3.060586
232 {M. 1870. R. High,M. 4275. R. High,M. 647. R. High} => {M. 646. R. High} 0.1085 0.8238421 3.059198
117 {M. 1817. R. High,M. 647. R. High} => {M. 646. R. High} 0.1026 0.8234350 3.057687
```

From the above figure we can interpret that, If the user rates Movies 2526 (The Fugitive), 2749 (The hunt for red october) and 647 (Terminator 2: Judgement Day) as high then he will also prefer movie 646 (The Terminator) and rate it 'high'.

4.c.

lift: The strength of a rule is indicated by lift. It is indicated over a co-occurrence of antecedent and consequent. It gives the details of the improvement i.e increase in probability of the consequent for a given antecedent.

$$(\text{Rule Support}) / (\text{Support}(\text{Antecedent}) * \text{Support}(\text{Consequent}))$$

Use subset command to list all the rules with lift > 3.0

```
inspect(sort(subset(apr, subset = lift > 3), by = "lift"))
```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> inspect(head(sort(subset(apr, subset = lift > 3.0), by = "lift"),50))

```

	lhs	rhs	support	confidence	lift
199	{M. 2250. R. High, M. 2936. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1025	0.8464079	3.142993
205	{M. 2526. R. High, M. 2749. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1007	0.8440905	3.134387
215	{M. 2250. R. High, M. 4275. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1157	0.8390138	3.115536
224	{M. 2526. R. High, M. 4275. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1119	0.8369484	3.107866
208	{M. 2250. R. High, M. 2526. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1158	0.8324946	3.091328
203	{M. 2250. R. High, M. 2749. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1006	0.8293487	3.079646
235	{M. 4275. R. High, M. 4712. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1112	0.8261516	3.067774
218	{M. 2250. R. High, M. 4712. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1130	0.8242159	3.060586
232	{M. 1870. R. High, M. 4275. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1085	0.8238421	3.059198
117	{M. 1817. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1026	0.8234350	3.057687
225	{M. 2526. R. High, M. 4712. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1075	0.8231240	3.056532
220	{M. 2526. R. High, M. 5407. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1012	0.8214286	3.050236
222	{M. 1870. R. High, M. 2526. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1072	0.8195719	3.043341
135	{M. 2936. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1164	0.8185654	3.039604
212	{M. 1870. R. High, M. 2250. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1084	0.8181132	3.037925
210	{M. 2250. R. High, M. 5407. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1038	0.8166798	3.032602
228	{M. 4275. R. High, M. 5407. R. High, M. 647. R. High}	=> {M. 646. R. High}	0.1066	0.8149847	3.026308

From the above figure we can interpret that, If the user rates Movies 2250 (Die Hard), 2936 (Lethal Weapon) and 647 (Terminator 2: Judgement Day) as high then he will also prefer movie 646 (The Terminator) and rate it 'high'.

5.

```
require(recommenderlab)
```

5.a.

```
ratings <- scan("ratings.txt", what="list", sep = "|")
```

```
ratings.matrix <- matrix(ratings, ncol=3, byrow=T)
```

```
class(ratings.matrix) <- "numeric"
```

```
ratings.sparseMatrix <- sparseMatrix(i=ratings.matrix[,1], j=ratings.matrix[,2], x=ratings.matrix[,3])
```

```
dimnames(ratings.sparseMatrix) <- list(user=paste("U", 1:10000), Movie=paste("M.", 1:7223))
```

```
dim(ratings.sparseMatrix)
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/ <R>
> head(ratings.sparseMatrix)
6 x 7223 sparse matrix of class "dgCMatrix"
[[ suppressing 76 column names 'M. 1', 'M. 2', 'M. 3' ... ]]

R 1 .....
R 2 .....
R 3 .....
R 4 .....
R 5 .....
R 6 .....

.....suppressing columns in show(); maybe adjust 'options(max.print=*)'
> dim(ratings.sparseMatrix)
[1] 10000 7223
```

Dimensions of the sparse matrix are: 10000 7223

5.b

```
realRatingMatrix <- new("realRatingMatrix", data=ratings.sparseMatrix)
```

```
realRatingMatrix.split <- floor(nrow(realRatingMatrix)*0.8)
```

```
realRatingMatrix.split.sampled <- sample(nrow(realRatingMatrix), realRatingMatrix.split)
```

```
ratingtrain <- realRatingMatrix[realRatingMatrix.split.sampled,]
```

```
recommend <- Recommender(ratingtrain, method="UBCF")
```

```
predict <- predict(recommend, realRatingMatrix[10000], n=10)
```

```
predict
```

```
as(predict, "list")
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/ <R>
> recommend <- Recommender(ratingtrain, method="UBCF")
> predict <- predict(recommend, realRatingMatrix[10000], n=5)
> predict
Recommendations as 'topNList' with n = 5 for 1 users.
> as(predict, "list")
[[1]]
[1] "M. 2242" "M. 3084" "M. 2434" "M. 2584" "M. 3774"
```

Top 5 Movie recommendations for user #10000 are:

2242 : Crying Game

3084 : Mission Impossible

2434 : Fargo

2584 : Gone with the wind

3774 : Rain man

5.c.

```
predict500 <- predict(recommend, realRatingMatrix[500], n=1)
```

```
predict500
```

```
as(predict500,"list")
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Movies_data/
> predict500<-predict(recommend,realRatingMatrix[500],n=1)
> predict500
Recommendations as 'topNList' with n = 1 for 1 users.
> as(predict500,"list")
[[1]]
[1] "M. 4349"
```

Highest predicted rating movie for the user #500:

4349 : Good will hunting

6.

```
require(tm)
```

6.a.

```
autosData <- DirSource(directory = ".")
news.corpus <- Corpus(DirSource(directory = "."))
length(news.corpus)
```

folder: 'rec.autos'

```
Console D:/Courses/Pattern Recognition/assignments/HW
> length(news.corpus)
[1] 990
> |
```

folder: 'rec.motorcycles'

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data
> news.corpus <- Corpus(DirSource(directory = "."))
> length(news.corpus)
[1] 996
```

To print the corpus entry corresponding to rec.autos/103806:

```
ds<- DirSource("D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos")
news.corpus<-Corpus(ds, readerControl=list(language="eng", reader=readPlain))
```

```
for(i in 1:length(news.corpus))
{
  if(names(b)[[i]]==103806){
    x<- i
    break
  }
}
```

```
print(i)
```

Output: 980

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> ds<- DirSource("D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos")
> news.corpus<-Corpus(ds, readerControl=list(language="eng", reader=readPlain))
>
> for(i in 1:length(news.corpus))
+ {
+   if(names(b)[[i]]==103806){
+     x<- i
+     break
+   }
+ }
> print(i)
[1] 980
```

6.b.

Initial file 103806

news.corpus[[980]]\$content # 103806 file is located in the location 980 in the corpus according to the previous solution.

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus[[980]]$content
[1] "From: cheekeen@tartarus.uwa.edu.au (Desmond Chan)"
[2] "Subject: Re: Honda clutch chatter"
[3] "Organization: The University of Western Australia"
[4] "Lines: 8"
[5] "NNTP-Posting-Host: tartarus.uwa.edu.au"
[6] "X-Newsreader: NN version 6.4.19 #1"
[7] ""
[8] "    I also experience this kinda problem in my 89 BMW 318. During cold"
[9] "start ups, the clutch seems to be sticky and everytime i drive out, for"
[10] "about 5km, the clutch seems to stick onto somewhere that if i depress"
[11] "the clutch, the whole chassis moves along. But after preheating, it"
[12] "becomes smooth again. I think that your suggestion of being some"
[13] "humidity is right but there should be some remedy. I also found out that"
[14] "my clutch is already thin but still alright for a couple grand more!"
[15] ""

```

Remove punctuation

```

news.corpus <- tm_map(news.corpus, content_transformer(removePunctuation))
news.corpus[[980]]$content

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus <- tm_map(news.corpus, content_transformer(removePunctuation))
> news.corpus[[980]]$content
[1] "From cheekeentartarusuwaeduau Desmond Chan"
[2] "Subject Re Honda clutch chatter"
[3] "Organization The University of Western Australia"
[4] "Lines 8"
[5] "NNTPPostingHost tartarusuwaeduau"
[6] "XNewsreader NN version 6419 1"
[7] ""
[8] "    I also experience this kinda problem in my 89 BMW 318 During cold"
[9] "start ups the clutch seems to be sticky and everytime i drive out for"
[10] "about 5km the clutch seems to stick onto somewhere that if i depress"
[11] "the clutch the whole chassis moves along But after preheating it"
[12] "becomes smooth again I think that your suggestion of being some"
[13] "humidity is right but there should be some remedy I also found out that"
[14] "my clutch is already thin but still alright for a couple grand more"
[15] ""

```

Remove Numbers:

```

news.corpus <- tm_map(news.corpus, content_transformer(removeNumbers))
news.corpus[[980]]$content

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus <- tm_map(news.corpus, content_transformer(removeNumbers))
> news.corpus[[980]]$content
[1] "From cheekeentartarusuwaeduau Desmond Chan"
[2] "Subject Re Honda clutch chatter"
[3] "Organization The University of Western Australia"
[4] "Lines "
[5] "NNTPPostingHost tartarusuwaeduau"
[6] "XNewsreader NN version "
[7] ""
[8] "    I also experience this kinda problem in my BMW During cold"
[9] "start ups the clutch seems to be sticky and everytime i drive out for"
[10] "about km the clutch seems to stick onto somewhere that if i depress"
[11] "the clutch the whole chassis moves along But after preheating it"
[12] "becomes smooth again I think that your suggestion of being some"
[13] "humidity is right but there should be some remedy I also found out that"
[14] "my clutch is already thin but still alright for a couple grand more"
[15] ""

```

Tolower:

```

news.corpus <- tm_map(news.corpus, content_transformer(tolower))
news.corpus[[980]]$content

```



```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus <- tm_map(news.corpus, content_transformer(tolower))
> news.corpus[[980]]$content
[1] "from cheekeentartarusuwaedua desmond chan"
[2] "subject re honda clutch chatter"
[3] "organization the university of western australia"
[4] "lines "
[5] "nntppostinghost tartarusuwaedua"
[6] "xnewsreader nn version "
[7] ""
[8] " i also experience this kinda problem in my bmw during cold"
[9] "start ups the clutch seems to be sticky and everytime i drive out for"
[10] "about km the clutch seems to stick onto somewhere that if i depress"
[11] "the clutch the whole chassis moves along but after preheating it"
[12] "becomes smooth again i think that your suggestion of being some"
[13] "humudity is right but there should be some remedy i also found out that"
[14] "my clutch is already thin but still alright for a couple grand more"
[15] ""

```

removeWords:

```

news.corpus <- tm_map(news.corpus, removeWords, stopwords("english"))
news.corpus[[980]]$content

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus <- tm_map(news.corpus, removeWords, stopwords("english"))
> news.corpus[[980]]$content
[1] "cheekeentartarusuwaedua desmond chan" "subject re honda clutch chatter"
[3] "organization university western australia" "lines "
[5] "nntppostinghost tartarusuwaedua" "xnewsreader nn version "
[7] "" " also experience kinda problem bmw cold"
[9] "start ups clutch seems sticky everytime drive " " km clutch seems stick onto somewhere depress"
[11] " clutch whole chassis moves along preheating " "becomes smooth think suggestion "
[13] "humudity right remedy also found " " clutch already thin still alright couple grand "
[15] ""

```

stripWhitespace:

```

news.corpus <- tm_map(news.corpus, content_transformer(stripWhitespace))
news.corpus[[980]]$content

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> news.corpus <- tm_map(news.corpus, content_transformer(stripWhitespace))
> news.corpus[[980]]$content
[1] "cheekeentartarusuwaedua desmond chan" "subject re honda clutch chatter"
[3] "organization university western australia" "lines "
[5] "nntppostinghost tartarusuwaedua" "xnewsreader nn version "
[7] "" " also experience kinda problem bmw cold"
[9] "start ups clutch seems sticky everytime drive " " km clutch seems stick onto somewhere depress"
[11] " clutch whole chassis moves along preheating " "becomes smooth think suggestion "
[13] "humudity right remedy also found " " clutch already thin still alright couple grand "
[15] ""

```

```

news.corpus <- Corpus(VectorSource((news.corpus)))

```

6.c.

```

dtm <- DocumentTermMatrix(news.corpus, control = list(minWordLength = 1, minDocFreq = 1, weighting = function(x)
weightTfIdf(x, normalize = FALSE)))

```

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> dtm
<<DocumentTermMatrix (documents: 990, terms: 14229)>>
Non-/sparse entries: 88152/13998558
sparsity : 99%
Maximal term length: 157
weighting : term frequency - inverse document frequency (tf-idf)

```

dim(dtm)

```

Console D:/Courses/Pattern Recognition/assignments/HW4/data/Newsgroup_data/Newsgroup_data/rec.autos/
> dim(dtm)
[1] 990 14229

```

Dimensions are: 990 14229

6.d.

```
inspectWords = inspect(DocumentTermMatrix(news.corpus[980],list(dictionary=c("bmw","clutch","mother"))))
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/
> inspectWords = inspect(DocumentTermMatrix(news.corpus[980],list(dictionary=c("bmw","clutch","mother"))))
<<DocumentTermMatrix (documents: 1, terms: 3)>>
Non-/sparse entries: 2/1
Sparsity           : 33%
Maximal term length: 6
Weighting           : term frequency (tf)

      Terms
Docs   bmw clutch mother
103806  1     5     0
```

From the above it can be seen that the word 'bmw' is present 1 time, 'clutch' is present 5 times and the word 'mother' is present 0 times in the file number 103806.

The results match with the expected outputs.

7

7.a.

```
azdata=read.table("az-5000.txt",header=T)
training<-sample(1:5000,4000)
trainingData<-azdata[training,]
aztestData<-azdata[-training,]
priors<-c(rep(1/26,26))
library(MASS)
azdatalda<-lda(char~,azdata,subset=training,prior=priors)
myprediction<-predict(azdatalda,newdata=azdata[-training,],type="response")
conformmat<-table(azdata[-training,]$char,myprediction$class)
conformmat
```

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/
> conformmat<-table(azdata[-training,]$char,myprediction$class)
> conformmat
```

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
a	28	0	0	7	0	0	1	0	0	0	0	0	0	0	0	0	1	0	0	1	0	0	0	3	1	3
b	0	33	0	0	2	1	0	2	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
c	0	0	34	0	3	0	0	2	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
d	0	0	0	20	0	0	0	0	0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	6	0	0
e	0	0	0	0	36	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
f	0	0	0	1	0	27	0	0	0	0	0	3	0	0	0	5	0	3	0	5	0	0	0	1	1	0
g	0	0	0	0	0	1	25	0	0	0	0	0	0	0	0	0	3	0	5	0	0	0	0	1	2	0
h	0	1	0	0	1	1	0	21	0	0	2	0	0	3	0	0	0	0	0	0	1	0	1	0	0	0
i	0	0	0	0	0	1	0	0	27	1	0	3	0	0	0	1	0	0	0	4	0	0	0	0	0	0
j	0	0	0	0	0	0	0	0	0	6	29	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
k	0	0	2	0	0	0	0	0	5	0	0	29	1	2	2	0	0	0	0	0	1	0	1	0	0	0
l	0	0	1	1	3	2	0	0	0	0	0	36	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	0	0	0	0	0	0	0	0	0	0	1	0	30	0	0	0	0	0	0	0	1	0	1	1	0	0
n	1	0	0	0	0	0	2	0	0	0	0	0	0	22	0	0	0	0	0	0	2	0	1	0	0	0
o	0	0	2	0	0	0	0	0	0	0	0	0	0	0	33	0	0	0	0	0	0	0	0	1	0	0
p	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	48	0	4	0	0	0	0	0	0	1	0
q	0	0	0	0	1	0	5	0	0	0	0	1	0	0	0	0	27	0	1	0	0	0	0	0	2	0
r	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	45	0	1	1	2	0	0	1	2	0
s	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	40	0	0	0	0	0	0	0	0
t	0	1	0	0	0	1	0	0	1	0	0	1	0	0	0	0	2	0	27	0	0	0	1	0	0	0
u	2	0	0	0	0	0	2	0	0	0	0	1	2	1	0	0	0	0	0	23	3	0	0	0	0	0
v	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0	1	0	0	3	36	0	0	0	0
w	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	32	0	0	0	0
x	0	1	1	3	0	0	0	0	0	0	2	0	0	0	0	1	1	0	1	2	2	0	20	6	0	0
y	0	0	0	0	0	2	2	0	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	35	0	0
z	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0	0	0	0	28	0

Making the diagonal '0'


```
for(k in 1:26){
  conformmat[k,k]=0
}
```

Conformmat

```
Console D:/Courses/Pattern Recognition/assignments/HW4/data/
> conformmat
```

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
a	0	0	0	7	0	0	1	0	0	0	0	0	0	0	0	1	0	0	1	0	0	0	3	1	3	0
b	0	0	0	0	2	1	0	2	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
c	0	0	0	0	3	0	0	2	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
d	0	0	0	0	0	0	0	0	0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	6	0	0
e	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
f	0	0	0	1	0	0	0	0	0	0	0	3	0	0	0	5	0	3	0	5	0	0	0	1	1	0
g	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	3	0	5	0	0	0	0	1	2	0	0
h	0	1	0	0	1	1	0	0	0	2	0	0	3	0	0	0	0	0	0	1	0	1	0	0	0	0
i	0	0	0	0	1	0	0	0	1	0	3	0	0	0	1	0	0	0	4	0	0	0	0	0	0	0
j	0	0	0	0	0	0	0	6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
k	0	0	2	0	0	0	0	5	0	0	0	1	2	2	0	0	0	0	0	1	0	1	0	0	0	0
l	0	0	1	1	3	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0	1	1	0	0	0
n	1	0	0	0	0	0	2	0	0	0	0	0	0	0	0	0	0	0	0	2	0	1	0	0	0	0
o	0	0	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
p	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	48	0	4	0	0	0	0	0	0	1	0
q	0	0	0	0	1	0	5	0	0	0	1	0	0	0	0	0	27	0	1	0	0	0	0	0	2	0
r	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	45	0	1	1	2	0	0	1	2	0
s	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	40	0	0	0	0	0	0	0	0
t	0	1	0	0	0	1	0	0	1	0	0	1	0	0	0	0	2	0	27	0	0	0	1	0	0	0
u	2	0	0	0	0	0	0	2	0	0	0	1	2	1	0	0	0	0	0	23	3	0	0	0	0	0
v	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0	1	0	0	3	36	0	0	0	0
w	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	32	0	0	0	0
x	0	1	1	3	0	0	0	0	0	2	0	0	0	0	0	1	1	0	1	2	2	0	20	6	0	0
y	0	0	0	0	0	2	2	0	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	35	0	0
z	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0	0	0	0	28	0

```
image(z=conformmat,zlim=c(1,10),col=rainbow(4), axes=FALSE)
axis(1, at = seq(0, 1, length=length(colnames(conformmat))), labels=colnames(conformmat))
axis(2, at = seq(0, 1, length=length(colnames(conformmat))), labels=colnames(conformmat))
```



7.b.
Color with most confusion is Blue.
Pairs:
{a,d}, {j,i},{d,x},{x,y}